(a) (3 Punkte) The Bayes error for classification is:

- ☐ the lowest error achievable by a linear classifier.
- ☐ the lowest error achievable by a nonlinear quadratic classifier
- ☑ the lowest error achievable by the best possible classifier
- ☐ the lowest error achievable by a classifier assuming Gaussian-generated distributions.

(b) (3 Punkte) Independent Component Analysis can be achieved by:

- ☐ Finding the directions in the input space that maximize the variance of the projected data.
- ☐ Applying a whitening procedure to the data and running PCA on the whitened data.
- ☑ Finding the directions in the input space that maximize the skewness of the projected data.
- ☐ Finding the directions in the input space that minimize the variance of the projected data.

(c) (3 Punkte) The K-means algorithm:

- ☐ Is a convex algorithm that can be used to cluster the data.
- ☑ Is a nonconvex algorithm that can be used to cluster the data.
- ☐ Is a kernelized version of the means algorithm where the kernel is Gaussian.
- ☐ Is a kernelized version of the means algorithm where the kernel can be arbitrary.

(d) (3 Punkte) A biased estimator is sometimes used to:

- ☐ Reduce the risk of underfitting the data.
- ☑ Reduce the estimation error for high-dimensional data.
- ☐ Make the estimation procedure more sensitive to the observed data.
- ☐ None of the above. We should always favor an unbiased estimator.

(e) (3 Punkte) Which is **False**? The Restricted Boltzmann machine is:

- ☑ A machine learning method that is based on error backpropagation.
- ☐ A machine learning method that can learn initial weights for a neural network.
- ☐ A machine learning method that estimates binary probability distributions for the data.
- ☐ A machine learning method that can learn global and local features in the data.

(a) (3 Punkte) Let $f_1(x) \dots f_N(x)$ be a set of descriminats for classification. The classification decision is given by $c* = \underset{i}{\text{argmax}}\ f_i(x)$ Which of the following sets would produce the same classification as the one above?

- ☐ $g_i(x) = (f_i(x))^2$
- ☐ $h_i(x) = log(1 + exp(f_i(x)))$
- ☐ None of them
- ☑ Both of them

(b) (3 Punkte) Consider a two-class classification problem. A sufficient condition for the Bayes optimal classifier to be linear is:

- ☐ The data generating distributions for both classes are equivalent except for the mean
- ☐ The data generating distributions for both classes are Gaussian
- ☐ The data generating distributions for both classes have the same covariance
- ☑ None of the above

(c) (3 Punkte) The Fisher linear discriminant finds the projection that:

- ☐ Maximizes the margin between the two data generating distribution
- ☐ Maximizes the margin between the mean of the two data generating distribution
- ☐ Maximizes the ratio between the within-class variance and the between-class variance
- ☑ Minimizes the ratio between the within-class variance and the between-class variance

(d) (3 Punkte) Which is **false**? Let k be a Gaussian Kernel. A Gram Matrix K associated to this Kernel always satisfies:

- ☐ $K = K^T$
- ☐ $KK^T = I$
- ☐ All Eigenvalues are non-negative
- ☑ $\forall u \in R^N : u^T K u \geq 0$

(e) (3 Punkte) Error Backpropagation is a technique to:

- ☑ Efficiently compute the error gradient in a multilayer neural network.
- ☐ Efficiently compute the error gradient in restricted Boltzmann machine.
- ☐ Efficiently compute the prediction error of a multilayer neural network.
- ☐ Efficiently compute the prediction error of a restricted Boltzmann machine.

What …. is not a discriminant function:
a) P(w_c|x)
b)P(x|w_c)*P(w_c)
c)P(w_c|x)*P(w_c)^-1
d)P(x|w_c)^2*P(w_c)^2

What is likely an overfitted estimator?
a) High bias model
b) High variance model
c) Low bias model
d) Low variance model

What does Fisher discriminant optimize?
a)..
b)..
c) Maximize ratio Within class variance to between class variance
d) Minimize ratio Within class variance to between class variance

What does the constant C stand for in SWM?
a) ability of the decision boundary to be out of the margin
b) number of points not being classified correctly
c)
d)

(a) The Bayes error is

- ☐ the lowest error of a linear classifier.
- ☐ the expected error of a random linear classifier.
- ■ the error of any nonlinear classifier.
- ☐ the error of a naive BAYES classifier .

(b) The Fisher linear discriminant find the projection $y = w^\mathsf{T} x$ of the data that maximises

- ☐ the margin between the two data generating distributions.
- ☐ the within-class variance divided by the between-class variance.
- ☐ the margin between the means of the data generating distributions.
- ■ the between-class variance divided by the within-class variance.

(c) A biased estimator is used to

- ■ make the estimator less affected by the sampling of the data.
- ☐ make the estimation procedure more sensitive to the sample data.
- ☐ reduce the risk of underfitting the data.
- ☐ None of the above, an unbiased estimator is always better.

(d) Let $x_1, \ldots, x_N \in \mathbb{R}^d$ be unlabelled observations. Consider a GAUSSIAN kernel and its GRAM matrix $K \in \mathbb{R}^{N \times N}$. Which is always true?

- ☐ $K^T K = I$.
- ☐ $K K^T = I$.
- ■ $\forall u \in \mathbb{R}^N \; uKu \geq 0$.
- ☐ $\forall u \in \mathbb{R}^N \; uKu \leq 0$.

1. Given two normal distributions $p(x|w_1) \sim \mathcal{N}(\mu_1, \Sigma_1)$ and $p(x|w_2) \sim \mathcal{N}(\mu_2, \Sigma_2)$ what is a *necessary* and *sufficient* condition for the optimal decision boundary to be linear? (5pts)

   (a) $\Sigma_1 = \Sigma_2$
   (b) $\Sigma_1 = \Sigma_2, P(w_1) = P(w_2)$
   (c) ...
   (d) ...

2. We have a classifier that decides the class $\text{argmax}_{w_i} f_i(x)$ for the input x. What is a suitable discriminant functions f_i? (5pts)

   (a) $\sqrt{p(x|w_i)P(w_i)}$
   (b) $\log\left(p(x|w_i) + P(w_i)\right)$
   (c) ...
   (d) ...

3. K-means is (5pts)

   (a) a non-convex algorithm used to cluster data
   (b) a kernelized version of the means algorithm
   (c) ...
   (d) ...

4. Error backpropagation gives (5pts)

   (a) the gradient of the error function
   (b) the optimal direction in parameter space
   (c) ...
   (d)

Which of the following is false: Assume a boosted classifier consists of weak hypotheses (aka. weak classifiers) that are each of them implemented by a threshold neuron. In that case the boosted classifier:

- ☐ can be viewed as a two-layer neural network.
- ☐ can be trained by error backpropagation instead of AdaBoost.
- ☐ can represent nonlinear decision boundaries.
- ☐ can represent non-smooth decision boundaries.

Which of the following is true: A Product of Experts:

- ☐ is an extension of a mixture model where each mixture element is allowed to be non-Gaussian.
- ☐ is an extension of a mixture model where each mixture element can be Gaussian with non-isotropic covariance.
- ☐ allows to learn more global features compared to a mixture model.
- ☐ allows to learn more local features compared to a mixture model.

Which of the following is false: Gaussian kernel ridge regression:

- ☐ is an extension of ridge regression to non-linear models.
- ☐ admits a closed-form solution when minimized for least squares.
- ☐ learns smooth non-linear functions.
- ☐ assumes that the input data is drawn from a Gaussian distribution.

Which of the following is true: In learning theory, the VC (Vapnik-Chervonenkis) bound:

- ☐ is an upper bound to the generalization error of a trained ML classifier of any complexity.
- ☐ is a lower bound to the generalization error of a trained ML classifier of any complexity.
- ☐ is an upper bound to the generalization error of a trained ML classifier of limited complexity.
- ☐ is a lower bound to the generalization error of a trained ML classifier of limited complexity.

Question 1

Incorrect

Mark 0.00 out of 5.00

Which of the following is **True**: A Gaussian Process (GP):

a. defines a multivariate Gaussian distribution over output variables, with covariance determined by input similarity.

b. defines a multivariate Gaussian distribution over input variables, with covariance determined by output similarity.

c. defines a multivariate distribution over output variables, with input drawn from a Gaussian distribution. ✖

d. defines a multivariate Gaussian distribution over input variables.

Your answer is incorrect.

The correct answer is:
defines a multivariate Gaussian distribution over output variables, with covariance determined by input similarity.

Question 2

Correct

Mark 5.00 out of 5.00

Which of the following is **True**: In learning theory, the VC (Vapnik-Chervonenkis) bound:

a. Is an upper-bound to the generalization error of a trained ML classifier of any complexity.

b. Is a lower-bound to the generalization error of a trained ML classifier of any complexity.

c. Is an upper-bound to the generalization error of a trained ML classifier of limited complexity. ✔

d. Is a lower-bound to the generalization error of a trained ML classifier of limited complexity.

Your answer is correct.

The correct answer is:
Is an upper-bound to the generalization error of a trained ML classifier of limited complexity.

Question 3

Correct

Mark 5.00 out of 5.00

Which of the following is **True**: k-means:

a. Is a supervised learning algorithm similar to k-nearest neighbors.

b. Has a convex objective and always converges to the global optimum.

c. Learns a solution that depends on the initialization.

d. Is a supervised learning algorithm for representation learning.

Your answer is correct.

The correct answer is:
Learns a solution that depends on the initialization.

Question 4

Incorrect

Mark 0.00 out of 5.00

Which of the following is **True**: A Product of Experts:

a. Is an extension of a mixture model where each mixture element is forced to be Gaussian.

b. Is an extension of a mixture model where each mixture element can be Gaussian with non-isotropic covariance.

c. Learns less local features than a mixture model.

d. Is an extension of a mixture model where each mixture element can be non-Gaussian with isotropic covariance.

Your answer is incorrect.

The correct answer is:
Learns less local features than a mixture model.

i) What is the Bayes error.

ii) Something about the fisher-discriminant

iii) When do you use a biased estimator.

iv) What is the k-means algorithm

Question 1

Incorrect

Mark 0.00 out of 5.00

Which of the following is **True**: Let k be a Mercer (PSD and symmetric) kernel and $\boldsymbol{x}_1, \ldots, \boldsymbol{x}_N$ be an unlabeled dataset. A Gram matrix K of size $N \times N$ associated to this kernel and dataset always satisfies:

a. $KK^\top = I$. ✖

b. $K^\top = K^{-1}$.

c. $\forall\, u \in \mathbb{R}^N : \; u^\top K u \geq 0$.

d. $\forall_{i=1}^N \forall_{j=1}^N : K_{ij} > 0$.

Your answer is incorrect.

The correct answer is:
$\forall\, u \in \mathbb{R}^N : \; u^\top K u \geq 0$.

Question 2

Incorrect

Mark 0.00 out of 5.00

Which of the following is **False** : PCA finds directions in input space for which:

a. The projection of non-centered data has maximum variance.
b. The projection centered data has maximum variance.
c. The projection of non-centered data has maximum sum-of-squares.
d. The projection centered data has maximum sum-of-squares. ✖

Your answer is incorrect.

The correct answer is:
The projection of non-centered data has maximum sum-of-squares.

Question 3

Correct

Mark 5.00 out of 5.00

Which of the following is **True**: In explainable machine learning, Shapley values:

a. can be computed in the order of a single forward/backward pass.
b. requires an exponential number of function evaluations to be computed. ✔
c. requires $O(d)$ function evaluations, where d is the number of input dimensions.
d. is a self-explainable model that must be trained alongside the actual model of interest.

Your answer is correct.

The correct answer is:
requires an exponential number of function evaluations to be computed.

Question 4

Correct

Mark 5.00 out of 5.00

Which of the following is **True**: Layer-wise relevance propagation (LRP) is a method for explainable AI that:

a. can be applied to any black-box machine learning model.
b. assumes that the machine learning model has a neural network (or computational graph) structure. ✔
c. requires $O(d)$ function evaluations, where d is the number of input dimensions, in order to produce an explanation.
d. can be applied to any black-box model, with the only condition that the gradient w.r.t. the input features can be computed.

Your answer is correct.

The correct answer is:
assumes that the machine learning model has a neural network (or computational graph) structure.

Question 1
Incorrect
Mark 0.00 out of 5.00
Flag question

Which of the following is **True**: Let k be a Mercer (PSD and symmetric) kernel and $\boldsymbol{x}_1, \ldots, \boldsymbol{x}_N$ be an unlabeled dataset. A Gram matrix K of size $N \times N$ associated to this kernel and dataset always satisfies:

a. $KK^\top = I$.
b. $K^\top = K^{-1}$.
c. $\forall\, u \in \mathbb{R}^N : u^\top K u \geq 0$.
d. $\forall_{i=1}^N \forall_{j=1}^N : K_{ij} > 0$. ✖

Your answer is incorrect.
The correct answer is:
$\forall\, u \in \mathbb{R}^N : u^\top K u \geq 0$.

Question 2
Correct
Mark 5.00 out of 5.00
Flag question

Which of the following is **True**: Layer-wise relevance propagation (LRP) is a method for explainable AI that:

a. can be applied to any black-box machine learning model.
b. assumes that the machine learning model has a neural network (or computational graph) structure. ✔
c. requires $O(d)$ function evaluations, where d is the number of input dimensions, in order to produce an explanation.
d. can be applied to any black-box model, with the only condition that the gradient w.r.t. the input features can be computed.

Your answer is correct.
The correct answer is:
assumes that the machine learning model has a neural network (or computational graph) structure.

Question 3
Correct
Mark 5.00 out of 5.00
Flag question

Which of the following is **True**: A Random Forest

a. Is a collection of decision trees, where the decision is taken by averaging (or majority voting of) the output of the individual decision trees. ✔
b. Is a boosted classifier that makes use of decision trees as weak learners.
c. Is a special type of neural network architecture with restricted connectivity.
d. Is a special type of neural network architecture where connections between neurons are set randomly with some probability.

Your answer is correct.
The correct answer is:
Is a collection of decision trees, where the decision is taken by averaging (or majority voting of) the output of the individual decision trees.

Question 4
Correct
Mark 5.00 out of 5.00
Flag question

Which of the following is **True**: A Mixture Model:

a. Is a class of unsupervised models designed in a way that the objective function is always convex.
b. Should be used if fitting a single distribution (e.g. Gaussian) for maximum likelihood takes too long or does not converge.
c. Is a class of unsupervised models that approximates the probability density function from which the data is generated. ✔
d. Is a class of unsupervised models that can represent efficiently any probability density function.

Your answer is correct.
The correct answer is:
Is a class of unsupervised models that approximates the probability density function from which the data is generated.